

Neural Architecture Search for Joint Big/Little Dynamic Inference Networks

Kristoffer Petersen, Oliver Larsen, Oliver Svendsen, Aleksander Korsholm
University of Southern Denmark (SDU)

Supervisors: Emil Jørgensen Njor, Francesco Daghero

Abstract—Edge and battery-powered systems increasingly run deep neural networks on-device, where strict memory, latency, and energy budgets constrain how much computation each input is allowed. Reducing the average cost by adapting computation to input difficulty: a small “little” network handles easy inputs, and harder inputs are deferred to a larger “big” network. Designing such a big/little cascade is a coupled problem. The two networks must share a routing rule, a memory budget, and a calibration behaviour, yet most existing work either places multiple exits inside a single backbone or composes a cascade from networks trained independently. We ask whether the little and big networks should instead be searched *jointly* as a cascade pair under explicit microcontroller (MCU) deployment constraints. We propose a constrained NSGA-II co-search that selects little/big CNN pairs on CIFAR-10 and compare it against an independent NAS baseline. At a fixed routing threshold $\tau = 0.70$, the joint cascade reaches 80.69% accuracy at 2.08 M MACs against 74.03% at 2.22 M MACs for the independent baseline. Deployed on a 512 KiB microcontroller, the joint cascade further reduces measured runtime by 23% and route-aware energy by 22% over the same baseline.

I. INTRODUCTION

Deep neural networks have grown both more accurate and more expensive [?], [?]. The same models that deliver strong predictions in the cloud are increasingly being asked to run on-device, on microcontrollers and other battery-powered embedded systems where a typical flash and SRAM budget is under one megabyte and average inference energy is a first-class constraint [?], [?], [?]. Always-on classification, keyword spotting, and sensor processing in this “TinyML” regime has motivated benchmarks and tooling for on-device deep learning [?], but the gap between state-of-the-art network capacity and microcontroller resources remains large.

Dynamic inference narrows this gap by spending compute only where it is needed. Rather than running a fixed network on every input, a dynamic network adapts the inference path to each sample [?], [?]. A big/little cascade is one of the simplest such designs [?]: a cheap “little” network is evaluated on every input, and a more expressive “big” network is invoked only when the little network is uncertain. The cascade can decide this route with a learned gate [?], [?], a margin-based rule [?], or a confidence threshold [?], [?]. We use a maximum-softmax confidence threshold because it requires no extra routing model, has negligible runtime overhead, and is straightforward to deploy under embedded constraints. The full routing rule is given in Section II-A. Such a cascade reduces average compute only if the little network is accurate on easy

samples, the big network is strong on deferred samples, and the threshold routes examples correctly.

Designing such cascades is not just a matter of finding two accurate models. The two models must work together under a shared routing rule and a shared deployment budget. If the little model is too small, it may rarely exit correctly. If the big model receives too much capacity, the cascade may become accurate but expensive. If the confidence threshold is too low, confident wrong exits can damage accuracy; if it is too high, most samples defer and the cascade loses its computational benefit. These trade-offs make cascade design a multi-objective problem involving accuracy, MACs, exit ratio, memory, and calibration.

Neural architecture search (NAS) is a natural tool for this setting because it can explore many architectural design choices under explicit constraints [?]. Prior work has studied NAS for efficient models, multi-objective architecture search, and dynamic inference, including early-exit networks and model cascades [?], [?], [?]. However, many cascade systems either design models independently or focus on placing exits inside a single backbone (Table I). This leaves open the question of whether a little and big model should be searched jointly as a pair, so that the search objective directly reflects cascade behavior.

This report studies that question on CIFAR-10 under an edge-style deployment constraint. We compare a joint co-search method, in which NSGA-II searches complete little/big pairs against cascade accuracy and expected cascade MACs directly, against an independent baseline that searches the two roles separately as standalone classifiers and combines them afterwards. Both methods share the underlying CNN search space, the same *proxy-training* budget (a short 5-epoch run used only to rank candidates during search; final claims use full retraining), and the deployment constraints. The comparison therefore isolates the effect of *searching as a cascade* from the effect of architectural capacity.

Architecture search is only part of the problem, however, because cascade routing depends on confidence. Cross-entropy (CE) optimizes classification accuracy but does not reflect the fact that some samples should exit at the little model while others should defer, and miscalibrated confidence translates directly into harmful early exits [?], [?], [?]. We therefore retrain the selected joint architectures under four objectives: CE, CE with label smoothing, a cascade-aware co-training loss, and the cascade-aware loss combined with label smoothing. These

choices are not fully orthogonal: the cascade-aware loss is itself confidence-shaping and interacts with the architecture’s routing behaviour at the same threshold, so RQ2 is studied as an ablation over a fixed set of selected architectures rather than as an independent optimisation.

The report is organised around two research questions:

- **RQ1:** Does joint NAS co-search of a little/big model pair produce a better final cascade accuracy versus average cascade MACs trade-off than independently searched and combined model pairs under matched search spaces and edge-memory constraints?
- **RQ2:** After architectures are selected by joint NAS, does cascade-aware co-training improve cascade accuracy, routing quality, or calibration compared with standard CE and CE with label smoothing?

We also analyse how the operating threshold τ , the resulting exit ratio, and the little-model calibration jointly shape the deployed accuracy MACs trade-off. This is a follow-up analysis on the selected joint cascades rather than a separate research question.

The implementation and experiment scripts are available in the project repository.¹

Contributions. This work makes the following contributions:

- 1) A joint NSGA-II co-search procedure for big/little CNN cascades under a 450 KiB combined weight-plus-activation budget motivated by a 512 KiB MCU. The comparison keeps the search space, proxy budget, and deployment constraints matched between joint and independent NAS.
- 2) A cascade-aware co-training loss that replaces the hard routing decision with a soft exit probability based on the little network’s confidence, so that little-model and big-model cross-entropy are weighted by the route a sample is likely to take. Treated as an ablation (conditions C1-C4 in Section IV-E), it shows how retraining choices shift the same selected architectures between cheaper and more accurate operating points.
- 3) A microcontroller deployment study on a Raspberry Pi Pico 2 connecting the analytical accuracy-MACs frontier to measured runtime and route-aware energy while fitting within the TensorFlow Lite Micro [?] memory arenas of the target.
- 4) An analysis of how threshold τ , exit ratio ρ , calibration (ECE, Section II-C), and routing error jointly shape the deployed accuracy-MACs trade-off, explaining why a high exit ratio on the independent baseline does not translate into a stronger cascade.

The remainder of the report is structured as follows. Section II introduces the background concepts shared by the rest of the paper (threshold-routed cascades, neural architecture search, and calibration). Section III reviews dynamic inference, NAS for efficient models, and confidence calibration for cascade routing. Section IV describes the search space, the




Fig. 1. Threshold-based routing in the big/little cascade. Every sample pays the little-model cost F_l ; green denotes the early-exit path for samples with $c_l > \tau$, while orange denotes the deferred path that additionally pays the big-model cost F_b .

joint and independent NAS procedures, and the cascade-aware loss. Section V presents the experimental setup and evaluation metrics. Section VI reports the results for the two research questions and the supporting routing analysis. Section VII discusses the main findings, limitations, and future work, and Section VIII concludes the report.

II. BACKGROUND

This section defines the components shared by the rest of the paper: the big/little cascade routing rule, the architecture-search formulation, and the calibration metric used to interpret confidence-based routing.

A. Big/Little Cascade Routing

A big/little cascade chains a lightweight *little* model with a larger *big* model that acts as a fallback for samples the little model is uncertain about [?]. For each input x , the little model is evaluated first and produces logits z_l . We use the *max-softmax* confidence,

$$c_l = \max(\text{softmax}(z_l)), \quad (1)$$

because it requires no auxiliary gating network, adds negligible runtime overhead, and is straightforward to evaluate on embedded targets [?], [?]. The cascade applies a fixed threshold τ : if $c_l > \tau$, the little model’s prediction $\hat{y}_l$ is accepted and the cascade exits. Otherwise the input is deferred to the big model, which produces logits z_b and prediction $\hat{y}_b$. This threshold rule is illustrated in Figure 1.

Let $\rho \in [0, 1]$ denote the *exit ratio*, the fraction of samples handled by the little model alone. The expected per-sample compute cost (measured in multiply-accumulate operations, MACs) is then

$$F_{\text{cascade}} = F_l + (1 - \rho) F_b, \quad (2)$$

where F_l and F_b are the analytical MAC counts of the little and big models. The first term is paid on every input; only the deferred fraction $1 - \rho$ pays the big-model cost [?], [?].

B. Neural Architecture Search

Neural Architecture Search (NAS) refers to a class of *algorithms* that automate architecture design [?]. A NAS algorithm is defined by three components: a *search space* of candidate architectures, a *search strategy* that explores it, and a *performance-estimation method* that scores candidates.

¹Neural-Efficiency-Research-Unit-NERU/nas-dynamic-big-little-cascade

Modern NAS methods range from reinforcement-learning controllers [?], through evolutionary algorithms [?], [?], to differentiable relaxations [?] and weight-sharing supernetworks [?]. Because full retraining of every candidate is computationally prohibitive, practical NAS uses a proxy: a short training run (here, 5 epochs) whose final accuracy is used as a ranking signal for the search, with longer retraining applied only to the selected architectures [?], [?]. NAS additionally optimizes non-accuracy targets such as MACs, memory, or latency [?], [?], [?]; In this report, we use NSGA-II, a multi-objective algorithm[?], because it can evaluate architectures using more than one objective at the same time, such as accuracy and compute cost. This makes it suitable for mixed discrete architecture variables and produces an explicit Pareto front for the accuracy-compute trade-off.

C. Confidence Calibration and Expected Calibration Error

A classifier is *well calibrated* when its predicted confidence matches its empirical accuracy: among samples with $c_i \approx 0.7$, the model should be correct about 70% of the time. Modern deep networks are typically overconfident [?], which is consequential for cascades because miscalibrated confidence directly biases the threshold router toward unsafe early exits [?]. We quantify calibration using Expected Calibration Error (ECE) [?], [?]. Test samples are sorted by predicted confidence and grouped into K equal-width bins $B_1, \dots, B_K$; ECE is the weighted gap between the average confidence and the average accuracy in each bin,

$$\text{ECE} = \sum_{k=1}^K \frac{|B_k|}{N} |\text{acc}(B_k) - \text{conf}(B_k)|. \quad (3)$$

We use $K = 15$ bins, following common practice [?], and report ECE for the little model since its confidence is the quantity the threshold router actually consumes.

III. RELATED WORK

A. Dynamic Inference and Big/Little Cascading

Dynamic neural networks adapt their computation to each input instead of using a fixed inference path for all samples. Han et al. survey this area and group dynamic models into instance-wise, spatial-wise, and temporal-wise adaptation methods [?]. Big/little cascades are an instance-wise dynamic inference method: an inexpensive model handles easy samples, while difficult samples are routed to a more expensive model. Park et al. proposed an early big/little DNN system for low-power inference, showing that a small network can filter easy inputs before invoking a larger network [?]. This idea is attractive for embedded systems because the expected cost depends on the exit ratio, not only on the worst-case cost of the largest model.

Closely related work studies early-exit networks, where auxiliary classifiers are attached to intermediate layers of a single backbone. BranchyNet introduced side branches that allow confident samples to exit before reaching the final layer [?]. MSDNet extends this idea with multi-scale feature

representations for anytime and budgeted classification [?]. Shallow-Deep Networks analyze “overthinking”, where a deep network can change an already-correct intermediate prediction into an incorrect final one, and use early exits to reduce unnecessary computation [?]. These methods demonstrate the value of adaptive computation, but they primarily modify a single backbone. Our work instead studies a cascade of two independent CNNs and asks how capacity should be allocated between the little and big models.

The routing rule is central to both early-exit networks and model cascades. Most systems use confidence thresholds: a sample exits if the intermediate or little model is sufficiently confident. This makes calibration important, because overconfident wrong predictions directly become wrong early exits. Learning to Cascade addresses this issue by calibrating confidence specifically for cascade systems, showing that standard calibration methods do not always optimize the accuracy-cost trade-off of a cascade [?]. This motivates our separate analysis of confidence, routing errors, and retraining losses.

B. NAS for Dynamic Inference

NAS automates architecture design using a defined search space, a search strategy, and a performance-estimation method [?]. It is particularly relevant for dynamic inference because accuracy and compute must be optimized jointly. Multi-objective evolutionary NAS is a common fit for this setting, since it can produce a Pareto front rather than a single scalarized solution. NSGA-II is a standard elitist multi-objective genetic algorithm [?], and NSGA-Net applies multi-objective evolutionary search to neural architectures [?]. We follow this general multi-objective view, but apply it to cascade pairs under explicit memory and parameter constraints.

Several recent methods apply NAS to early-exit neural networks. EDANAS searches early-exit networks by adapting an Once-for-All style search space and jointly considering accuracy and MACs [?]. NACHOS extends this line toward hardware-constrained early-exit design, searching for networks that satisfy accuracy and MAC constraints at inference time [?]. These approaches are close in spirit because they use architecture search for adaptive inference. However, they focus on adding or optimizing exits within a single network. Our method searches two separate role-specific networks and optimizes their behavior as a cascade pair.

Other work focuses on learning the exit decision itself. JEI-DNN jointly learns the gating mechanism and intermediate inference modules, addressing the train-test mismatch that arises when confidence-threshold gates are trained separately from the prediction modules [?]. This line of work is complementary to ours: JEI-DNN improves learned gating inside dynamic neural networks, while our study keeps a simple threshold router and asks whether the architectures of the two cascade members should be searched jointly.

ENCAS is the closest related work on NAS for model cascades. It performs evolutionary neural cascade search across pretrained supernetworks to obtain a front of cascades trading accuracy against MACs [?]. Our work differs in two ways.

First, we co-search the little and big models from a small role-specific CNN search space under a strict edge-memory budget, rather than composing subnetworks from pretrained supernet-networks. Second, we explicitly compare joint co-search against an independently searched baseline and then study whether cascade-aware retraining improves the selected architectures.

C. Confidence Calibration for Cascade Routing

Modern neural networks are often miscalibrated: their predicted probabilities can be more confident than their empirical accuracy [?]. This is especially important for cascades because the confidence score is not only a diagnostic quantity, since it controls computation. A poorly calibrated little model can exit on samples that should have been deferred to the big model. As a result, optimizing standalone classification accuracy is not necessarily enough to optimize cascade accuracy.

Label smoothing is one common training modification that can affect both accuracy and calibration. Muller et al. show that label smoothing changes the structure of learned representations and can improve generalization, although its effect depends on the task and downstream use [?]. Focal loss has also been studied as a calibration-aware alternative; Mukhoti et al. show that it can improve calibration in some settings [?]. These results motivate treating the final training objective as part of the cascade design, rather than evaluating only the searched architecture.

to model accuracy [?]. This distinction matters because confidence is also the routing signal: an overconfident little model can save compute but produce an incorrect early exit, while an underconfident little model may defer too often and lose the computational benefit of the cascade. A useful confidence score should therefore separate samples the little model can handle from samples that are worth sending to the big model. Our cascade-aware loss follows the same motivation but is used as a retraining objective for a jointly searched big/little pair. It introduces soft routing weights during training so that little-model and big-model losses are weighted according to the expected route of each sample. The goal is to improve cascade behavior without replacing the simple threshold rule used at inference.

D. Resource-Constrained NAS

Hardware-aware NAS extends architecture search with deployment constraints such as latency, memory, energy, MACs, or parameter count. This is important because the most accurate architecture under an unconstrained search is often not deployable on edge hardware. Surveys of hardware-aware NAS describe how search objectives can be coupled to device-specific measurements or analytical proxies [?]. NetAdapt iteratively adapts a pretrained network to meet platform constraints [?], while ProxylessNAS searches directly on the target task and hardware instead of relying only on proxy tasks [?]. MobileNetV3 combines platform-aware search with efficient mobile building blocks [?].

During NAS, our search uses analytical resource constraints rather than measured latency: deployment memory, parameter

count, and average cascade MACs. This keeps architecture selection reproducible and independent of measurement noise from a specific hardware run. Measured Pico 2 latency is treated separately as a deployment evaluation after model selection. The memory constraint is applied to the whole cascade, not just to each model separately. This matters because a good cascade is a capacity-allocation problem: the little model must be large enough to exit correctly for the majority of easy samples, while the big model must still have enough capacity to improve deferred samples.

E. Gap in Existing Work

Table I summarizes how the most relevant methods relate to this work. Existing work strongly supports the usefulness of dynamic inference, early exits, confidence-aware routing, and hardware-aware NAS.

Before pointing at the gaps, it is worth noting that two distinct dynamic designs dominate the literature. *Single-backbone early-exit networks* (BranchyNet [?], MSDNet [?], EDANAS [?], NACHOS [?]) attach auxiliary classifiers at intermediate depths of one network and exit when an auxiliary head is confident; the backbone, exit positions, and exit heads are co-optimised. *Two-CNN cascades* (Park et al. [?], ENCAS [?]) instead use two separately-trained networks of different capacity and route between them by a confidence threshold. The two designs trade off different things. Early-exit backbones share feature extraction and so are typically more parameter-efficient at iso-capacity, but the same shared backbone also forces the cheap and expensive paths to live on the same memory arena and to be optimised against a single loss surface, which can be brittle on embedded targets. Two-CNN cascades pay the memory overhead of a second model, but gain a clean memory partition (each model has its own arena allocation, which matters on TensorFlow Lite Micro [?]), simpler training, and the ability to allocate capacity asymmetrically between an easy-case classifier and a hard-case fallback. Our work focuses on this two-CNN regime and asks whether the two role-specific architectures should additionally be *co-searched*, not just co-trained or co-routed.

Two gaps remain in this setting. First, most NAS methods for dynamic inference search early exits inside one backbone or assemble cascades from pretrained models; few study co-search of two independent role-specific architectures under a shared cascade objective. Second, the training objective for cascades is still not settled. A loss that improves routing may also harm branch accuracy, so it must be evaluated separately from the architecture search itself.

IV. METHODOLOGY

This section describes the architecture search space and the two design procedures compared in the experiments: joint NSGA-II co-search of complete big/little cascade pairs, and an independent baseline that searches the two roles separately. Both procedures evaluate candidates under the routing rule and cost model defined in Section II-A (Eq. 1 and Eq. 2). The

TABLE I
COMPARISON OF RELATED DYNAMIC-INFERENCING AND NAS METHODS.

Method	Main idea	Architecture setting	Relation to this work
BranchyNet [?]	Early exits with side classifiers, gated by max-softmax confidence	Single backbone	Source of our max-softmax routing rule, but inside one backbone
MSDNet [?]	Multi-scale anytime/budgeted classification	Single backbone	Shows adaptive compute improves efficiency
Park et al. [?]	Big/little low-power inference; samples routed by little-model softmax confidence	Two manually-designed models	Foundational cascade design we follow, but architectures are not searched
EDANAS [?]	NAS for early-exit networks	Single early-exit network	Searches exits, not separate little/big models
NACHOS [?]	Hardware-constrained early-exit NAS	Single early-exit network	Adds hardware constraints to early-exit NAS
JEI-DNN [?]	Jointly learned exit and inference modules	Learned gating in dynamic network	Learns gating; does not co-search cascade pair
ENCAS [?]	Evolutionary search over model cascades	Pretrained supernetwork cascades	Closest cascade NAS; different search setting
Ours	Joint NSGA-II co-search of little/big CNNs	Two independent role-specific models	Compares joint vs. independent search and losses

TABLE II
CNN SEARCH SPACE USED FOR JOINT AND INDEPENDENT NAS.

Component	Little model	Big model
Blocks	1-3	2-5
Channels	{8, 12, 16, 24, 32}	{16, 24, 32, 48, 64}
Layers/block	{1, 2}	
Kernel size	{3, 5}	
Convolution	standard or depthwise separable	
Residual	enabled or disabled	

TABLE III
SEARCH CONSTRAINTS APPLIED TO EVERY CANDIDATE. THE 450 KiB COMBINED CAP LEAVES ≈ 60 KiB OF HEADROOM AGAINST THE 512 KiB TARGET FOR STACK, FRAMEWORK CODE, AND RUNTIME BUFFERS.

Constraint	Value
Combined weight+activation budget	≤ 450 KiB
Combined parameter window	50K-100K
Little-model parameter range	15K-30K
Big-model parameter range	40K-85K
Role-order (parameters)	$ \theta_b \geq \theta_l $
Role-order (blocks)	$\text{blocks}(b) > \text{blocks}(l)$

cascade-aware retraining loss used for the RQ2 ablation is also introduced here.

Figure 3 contrasts the joint and independent pipelines side by side. The joint pipeline (steps J1-J7) treats the cascade as a coupled object as soon as proxy training begins, while the independent baseline (steps I1-I7) does not assemble a cascade until I4; the rest of this section refers to these step labels.

A. Search Space and Resource Constraints

Both models are sampled from the same CNN search space, but with role-specific channel ranges and depths. The little model uses the smaller role configuration, while the big model uses the larger one. Table II summarizes the searchable block counts, channel choices, and per-block operations used by both the joint and independent NAS experiments.

The search space is intentionally small and discrete. This keeps evolutionary search feasible under the available compute budget while still allowing meaningful variation in depth, width, convolution type, and residual connections. The little and big channel ranges are separated so that the little model is biased toward cheaper architectures, while the big model has enough capacity to serve as a fallback for harder samples. The block ranges follow the same role separation: the little model is limited to fewer stages, while the big model can use deeper feature extraction.

All candidates must satisfy a set of deployment constraints motivated by fitting the cascade on a 512 KiB-class microcontroller. The deployment-memory estimate sums both models' float32 weights, the input buffer, and the larger of the two peak activation footprints (the models are evaluated sequentially, so their activation arenas can be re-used). The constraints are summarised in Table III and apply identically to the joint and independent NAS procedures.

The lower bounds prevent the search from exploiting models that are cheap but trivially weak. The upper bounds keep the two roles within a shared capacity envelope. The big-model upper bound (85K) is higher than the little-model upper bound (30K) so that a 15K little model can still be paired with an 85K big model without exceeding the 100K combined cap. The role-order constraints prevent the first-stage model from being larger than the fallback. All memory and MAC numbers in this section are *analytical* estimates computed in float32: direct Pico 2 latency and TensorFlow Lite Micro [?] arena sizes are reported separately in Section VI-D after the architectures have been exported.

B. Joint Co-Search with NSGA-II

The joint NAS experiment searches little/big model pairs directly. Each candidate (J1 in Figure 3) encodes both archi-

tures and is first checked against the hard constraints of Table III (J2); infeasible candidates are rejected analytically and are not proxy-trained. The initial sampler and offspring repair operator keep drawing or repairing candidates until the population contains feasible pairs, so infeasible samples do not reduce the number of proxy-trained candidates in a generation. Feasible candidates proceed to J3, where the pair is proxy-trained for 5 epochs using the cascade-aware shared loss (Section IV-D), and J4, where the trained pair is evaluated as a cascade at the fixed threshold $\tau = 0.70$. The 5-epoch proxy is the same heuristic used by short-budget NAS methods such as regularised evolution [?] and DARTS [?]; long enough to separate clearly weak architectures from promising ones, short enough to evaluate hundreds of candidates within the available compute.

NSGA-II [?] (implemented via pymoo [?]) optimises two objectives simultaneously: *maximise* proxy cascade accuracy and *minimise* expected cascade MACs (Equation. 2). We do not collapse the two objectives into a scalar reward because the relative weighting between accuracy and compute is the engineering decision an embedded deployment should make post hoc, not an arbitrary search-time constant. Candidates whose exit ratio ρ falls outside the routing window $[0.20, 0.80]$ receive a soft penalty of 0.20 on the accuracy objective. This discourages degenerate solutions that always defer to the big model (which would dominate the accuracy objective at high MACs) or always exit at the little model (which would dominate the MACs objective at low accuracy). The threshold τ is held fixed during search to prevent the search from trivially raising it to make the cascade defer everything. We use $\tau = 0.70$ as the main operating point because it is permissive enough to produce meaningful early exits under the short proxy-training budget, while still filtering clearly low-confidence little-model predictions. Final operating points are swept post hoc. Figure 2 visualises the 400 proxy-trained candidates across the 20 generations of the joint search and shows how NSGA-II’s elitist selection concentrates later generations in the high-accuracy, low-MACs region of the feasible space.

After the search, three candidates are chosen for full retraining from the proxy Pareto region. Let $\mathcal{P}$ be the set of proxy-Pareto-optimal candidates and $F_b(\pi)$ be the big model’s standalone MACs for candidate $\pi \in \mathcal{P}$. The selector first filters

$$\mathcal{P}_{cs} = \{ \pi \in \mathcal{P} \mid F_{\text{cascade}}(\pi) < F_b(\pi) \}, \quad (4)$$

so that every selected pair is a genuine compute-saving cascade. Within $\mathcal{P}_{cs}$ the chosen candidates are the knee point and its two nearest *unique* neighbours, which captures the operating region where additional compute gives diminishing proxy-accuracy gains without selecting duplicate genotypes. The selection rule does not simply take the three highest proxy-accuracy candidates. The proxy frontier and the three selected candidates are shown in Figure 12 in Appendix A.

C. Independent NAS Baseline

The independent baseline tests whether joint pair search adds value beyond searching two good standalone models. Lit-




Fig. 2. Joint NSGA-II co-search progress across all 400 proxy-trained candidates (20 generations $\times$ 20 individuals). Each point is one little/big cascade pair; the x-axis is the average cascade MACs the search minimises and the y-axis is the proxy cascade accuracy it maximises. Colour encodes generation (darker = earlier). All plotted candidates satisfy the 450 KiB deployment-memory constraint, which is checked analytically before proxy training. Later generations cluster toward higher accuracy at moderate MACs as the elitist multi-objective selection moves the population toward the joint Pareto frontier.

tle and big architectures are optimised in *two* separate NSGA-II searches against the same CNN search space (Table II) and role-specific parameter ranges (Table III). Each search optimises standalone validation accuracy against parameter count, since cascade routing does not yet exist at this stage and expected cascade MACs cannot be computed. The two role-specific Pareto fronts are then combined *post hoc* by enumerating all little-Pareto $\times$ big-Pareto pairings, applying the pair-level feasibility checks (combined parameter budget, deployment-memory cap, role-order rules), and selecting the highest-scored feasible pairs. Enumerating after the role-specific searches, rather than fixing the best-of-role models, ensures that an 85K big model that cannot pair with a 30K little model is not discarded simply on those grounds: it can still be selected if it forms a valid cascade with a smaller little model.

The baseline is matched to the joint search wherever this is meaningful: it shares the CNN search space, the random seed, the proxy-training duration, the final threshold sweep, and the deployment constraints. It deliberately does not share the search objective, because the purpose of the baseline is to test whether independently strong models are sufficient. Figure 3 contrasts the two pipelines: joint NAS evaluates pairs as cascades during search, whereas the independent baseline delays cascade construction until after the two role-specific searches finish.

D. Cascade-Aware Loss

We use a cascade-aware retraining loss that keeps standard classification terms but reflects how the cascade routes samples at inference time. We use *co-training* in the sense of jointly optimising the little and big models with a shared loss, which




Fig. 3. Joint co-search pipeline (top, J1-J7) versus the independent NAS baseline (bottom, I1-I7). Stages are coloured by type: blue = sample / search, orange = feasibility filter, green = training (proxy or final), violet = evaluation, frontier, and selection. The joint pipeline sees the cascade as a coupled object from J3 onward, whereas the independent baseline does not assemble cascade pairs until I4. The text in Sections IV-B and IV-C refers to these step labels.




Fig. 4. Behavioral view of the cascade-aware loss. Green shows the high-confidence path, where samples are treated as likely early exits; orange shows the low-confidence path, where samples are treated as likely deferrals. The soft exit probability shifts the learning emphasis between the little and big models.

is unrelated to semi-supervised co-training. The hard routing decision (Eq. 1) is replaced during training by a differentiable soft exit probability

$$p_{\text{exit}} = \sigma(k(c_l - \tau)), \quad (5)$$

where σ is the sigmoid function and k is a routing sharpness parameter. Large k makes the soft decision approach the hard threshold rule, while small k smooths it out. We use $k = 20$, which gives $p_{\text{exit}} \approx 1$ for $c_l \gtrsim 0.85$ and $p_{\text{exit}} \approx 0$ for $c_l \lesssim 0.55$ at $\tau = 0.70$; this is sharp enough to behave like the inference-time threshold for most samples while keeping a small interval of gradient flow around τ .

The full cascade-aware loss has four terms:

$$\begin{aligned} \mathcal{L}_{\text{cascade}} = & \underbrace{(p_{\text{exit}} + \lambda_l) \text{CE}(z_l, y)}_{\text{little CE, routed}} \\ & + \underbrace{(1 - p_{\text{exit}} + \lambda_b) \text{CE}(z_b, y)}_{\text{big CE, routed}} \\ & + \underbrace{\lambda_d (1 - p_{\text{exit}})}_{\text{defer cost}} + \underbrace{\lambda_w p_{\text{exit}} \mathbf{1}[\hat{y}_l \neq y]}_{\text{wrong-exit penalty}}. \end{aligned} \quad (6)$$

The first two CE terms emphasise the little (resp. big) model on samples likely to exit (resp. defer); the constants λ_l and λ_b keep a small training signal flowing to both models even when the soft routing weight is near zero. The third term is a small defer cost that prevents the optimiser from sending every sample to the big model, which would maximise accuracy but eliminate the cascade’s compute saving. The fourth term penalises confident wrong exits, the failure mode that directly hurts a threshold-routed cascade. The routing-related terms are deliberately small relative to the CE terms. Large values dominate the classification objective and produce miscalibrated branches. Figure 4 illustrates these effects for high-confidence and low-confidence samples.

We use $\lambda_l = 0.75$, $\lambda_b = 0.05$, $\lambda_d = 0.0125$, and $\lambda_w = 0.075$ throughout. These values were selected by manual calibration on the CIFAR-10 validation split rather than by a full sweep: λ_l is the largest so that the little branch always receives a classification signal; λ_b is smaller because the big branch is mostly trained on deferred samples already; λ_d and λ_w are small enough to act as routing-aware regularisers without dominating the CE terms. A systematic sensitivity sweep is left to future work (§VII-F).

E. Retraining Loss Ablations

For RQ2, the three selected joint-NAS architectures are retrained from scratch under the four loss settings of Table IV. The ablation isolates whether improvements come from standard cross-entropy (C1), label smoothing (C2), the cascade-aware loss (C3), or its combination with label smoothing (C4). CE is the default baseline because it is the standard supervised classification objective [?], [?]. Label smoothing is included because cascade routing depends on confidence, and smoothing is a simple calibration-oriented baseline [?], [?].




Fig. 5. Class distribution of the first 5,000 CIFAR-10 training images used as the validation split.

TABLE IV
RETRAINING LOSS ABLATIONS.

Condition	Retraining loss
C1	Cross-entropy (CE)
C2	CE + label smoothing
C3	Cascade-aware loss
C4	Cascade-aware loss + label smoothing

V. EXPERIMENTAL SETUP

A. Datasets

All experiments use CIFAR-10 [?], a 10-class 32×32 image classification dataset with 50K training and 10K test images. We use a deterministic 45K/5K split of the training set: approximately 500 images per class (Figure 5) are held out for validation and the remaining 45K are used for training. Table IV;

The official 10K test split is kept untouched and used only for the final threshold-sweep evaluation. CIFAR-10 is small enough to support repeated NAS proxy training and ablation runs while remaining non-trivial for confidence-based routing and calibration analysis. Training uses random crop with 4-pixel padding and random horizontal flipping, followed by CIFAR-10 normalisation; validation and test use the same normalisation without augmentation.

B. Training Schedule

The training schedule is summarised in Table V. The joint search uses a population of 20 for 20 generations; the independent baseline uses 20 for 10 generations *per role*, since it runs one little-model and one big-model search sequentially. This gives the two methods approximately the same number of proxy-trained candidates (the quantity that dominates compute cost), even though the evolutionary depth per role differs. The independent baseline does not use the cascade-aware

TABLE V
TRAINING AND EVALUATION SCHEDULE USED IN THE EXPERIMENTS.

Stage	Setting	Value
Joint NAS	Population / generations	20 / 20
Independent NAS	Population / generations	20 / 10 per role
Proxy training	Epochs / batch size	5 / 128
Proxy training	Optimizer	SGD, lr 0.05
Joint C1-C4 retrain-ing	Epochs / batch size	50 / 128
Joint C1-C4 retrain-ing	Optimizer	SGD, lr 0.1
Joint C1-C4 retrain-ing	Scheduler	Cosine anneal-ing [?]
Independent final training	Epochs / learning rate	50 / 0.05
All final training	Momentum / weight decay	0.9 / $5 \cdot 10^{-4}$
Threshold sweep	Thresholds	0.50, 0.60, 0.70, 0.80, 0.90, 0.95

shared loss: it trains the selected little and big architectures as standalone CE classifiers, and the two are coupled only at evaluation via the threshold router.

Retraining starts from scratch under each of the four C1-C4 conditions in Table IV; proxy-trained weights are not re-used. The 50-epoch retraining budget is the longest schedule we could afford across all variants given the total compute budget. Many CIFAR-style setups train for 100-200 epochs; we discuss the implications of the shorter schedule in Section VII-F.

C. Evaluation Metrics

The primary metrics are cascade accuracy and expected cascade MACs (the cost model defined in Eq. 2). Secondary metrics are the exit ratio ρ , the standalone little- and big-model accuracies, parameter count, and the float32 deployment-memory estimate (both models' weights plus the input buffer plus the larger of the two peak activation footprints, since the two models are evaluated sequentially). All analytical metrics support NAS selection and the main accuracy-compute comparisons. Measured Pico 2 latency and TFLM arena sizes are reported separately (Section VI-D).

For routing and calibration we report the little-model ECE (Section II-C) and the *routing error rate*, the fraction of test samples on which the cascade exits at the little model with a wrong prediction even though the big model would have been correct. This captures the failure mode that directly hurts a threshold-routed cascade. Proxy-search and final-retrained results are reported separately: proxy metrics are 5-epoch validation accuracies used only for architecture selection, whereas final claims use 50-epoch retrained test-set threshold-sweep results.

D. Implementation Details

All experiments are implemented in PyTorch through a scripted pipeline that runs joint NAS, independent NAS,

retraining, metric aggregation, and plot generation. CUDA is used when available, CPU otherwise. All experiments use random seed 42, which is propagated to Python, NumPy, PyTorch, CUDA, the NSGA-II implementation, feasible-sampling and repair operators, and DataLoader shuffling. Experiment outputs are saved as CSV files, including search logs, Pareto fronts, selected architectures, training curves, and threshold sweeps. The reported results tables and figures are generated from these files, allowing the results to be recomputed without retraining.

E. Pico 2 Deployment Test Setup

After final retraining, the selected cascades were exported to ONNX and then to float32 TensorFlow Lite models for TensorFlow Lite Micro. The exported little and big models were converted into C arrays and compiled into Raspberry Pi Pico 2 firmware. The Pico firmware implements the same threshold router used in the offline evaluation: each image is first evaluated by the little model, the maximum softmax confidence is compared with $\tau = 0.70$, and only low-confidence samples are deferred to the big model.

Two Pico experiments were run. First, the normal latency and accuracy test used a USB serial protocol in which the host streamed CIFAR-10 test images to the Pico one sample at a time. This allowed the full 10,000-image CIFAR-10 test set to be evaluated without storing the entire dataset in Pico flash. For each sample, the Pico reported the predicted class, route decision, correctness, little-model time, big-model time when used, total cascade time, and analytical MAC counts for the selected path. This is the source of the deployment latency and route-split results reported in Table VIII.

Second, a route-aware energy test was run on a 1,000-image embedded CIFAR-10 subset. This version stores the test subset in Pico flash and runs without USB power during measurement. Figure 6 shows the measurement circuit. An Elegoo Uno R3 powers the Pico 2 through a $10\ \Omega$ high-side shunt connected to VSYS. The Uno samples the voltage before and after the shunt to estimate current and power, while Pico GP18 drives an inference trigger connected to Uno D2. In the route-aware build, Pico GP19 also indicates the route decision to Uno D3, with the high state representing deferred big-model execution. This setup separates idle time from the active inference window and records whether each measured inference used the little path or the deferred path.

VI. RESULTS

All reported results in this section use the final retrained checkpoints and the CIFAR-10 test split. Unless otherwise stated, results are reported at the fixed operating threshold $\tau = 0.70$, matching the main comparison used in the generated figures.

A. RQ1: Joint Co-Search vs. Independent NAS

Table VI summarizes the final joint and independent cascade results. The joint co-search pipeline occupies a stronger accuracy-MACs region than the independent baseline. C3 is

TABLE VI

RQ1 TEST-SET COMPARISON AT $\tau = 0.70$. VALUES ARE AVERAGED ACROSS THE THREE SELECTED RETRAINED CASCADES. C1-C4 SHARE THE SAME CO-SEARCHED ARCHITECTURES BUT DIFFER IN RETRAINING LOSS. THE INDEPENDENT BASELINE IS SEARCHED AND COMBINED SEPARATELY. ECE IS EXPECTED CALIBRATION ERROR (§II); LOWER IS BETTER.

Method	Acc. (%)	Exit ratio	MACs (M)	ECE
Joint NAS C1	79.64	0.55	1.69	0.0104
Joint NAS C2	80.33	0.43	2.01	0.0811
Joint NAS C3	79.69	0.55	1.67	0.0104
Joint NAS C4	80.69	0.40	2.08	0.0850
Independent NAS	74.03	0.57	2.22	0.0508

TABLE VII

RQ2 RETRAINING-LOSS ABLATION AT $\tau = 0.70$. VALUES ARE AVERAGED OVER THE THREE SELECTED JOINT ARCHITECTURES (A1, A2, A3).

Variant	Acc. (%)	Exit ratio	MACs (M)	Route err.
C1: CE	79.64	0.55	1.69	0.0231
C2: CE + smooth	80.33	0.43	2.01	0.0093
C3: Cascade loss	79.69	0.55	1.67	0.0242
C4: Cascade + smooth	80.69	0.40	2.08	0.0083

the cheaper cascade-aware variant, reaching 79.69% mean test accuracy at 1.67M average cascade MACs, while C4 reaches the highest mean accuracy, 80.69%, at a higher 2.08M average cascade MACs. The independent baseline reaches 74.03% at 2.22M average cascade MACs. Thus, at the fixed threshold used for the main comparison, the joint variants provide substantially higher accuracy, and the cheaper C3 operating point also uses less expected compute than the independent baseline.

The individual C4 cascades range from 79.77% to 81.71% accuracy at $\tau = 0.70$, while the independent cascades range from 72.67% to 75.96%. C3 gives a lower-compute joint operating point with similar accuracy to C1, whereas C4 gives the high-accuracy, high-deferral end of the ablation. This indicates that the joint pipeline is stronger than the independent baseline, but that the final operating point still depends on the retraining objective. The high-accuracy joint comparison is visualized in Figure 7.

B. RQ2: CE vs. Cascade-Aware Retraining

Table VII reports the retraining-loss ablation for the same three joint-NAS architectures. C1 and C3 compare CE against the cascade-aware loss without label smoothing. Their average cascade accuracies are close, but C3 slightly improves over C1 while also reducing expected MACs: C1 obtains 79.64% at 1.69M MACs, while C3 obtains 79.69% at 1.67M MACs. This is not a large standalone win, but it shows that the cascade-aware loss can move the operating point in the intended direction.

The smoothed variants, C2 and C4, achieve the highest average cascade accuracies. They also have substantially lower routing error rates than C1 and C3, but this comes with lower exit ratios and therefore higher expected cascade MACs. The ablation itself thus shows an accuracy-compute trade-off: C1




Fig. 6. Route-aware Pico 2 energy measurement setup using a 10 Ω high-side shunt, GP18 inference trigger, and GP19 route signal.

and C3 exit more often and are cheaper, while C2 and C4 defer more samples and obtain higher accuracy. C4 is the higher-accuracy, higher-compute variant; C3 is the cheaper cascade-aware alternative. These trends are shown in Figure 8.

C. Routing Behaviour and Threshold Sensitivity

The two research questions concern *which* architecture and *which* loss to use at a single operating threshold. This subsection complements them by examining *how* the selected joint cascades behave at the operating threshold and across a sweep of alternative thresholds.

At $\tau = 0.70$, the routing-error rate of the C4 variant is 0.0083, against 0.0231 and 0.0242 for C1 and C3 respectively (Table VII). The highest-accuracy variant also makes the fewest harmful early-exit decisions. ECE and routing error do not, however, move together: C1 and C3 have the lowest little-model ECE but the highest routing error. Global calibration is necessary but not sufficient to explain cascade quality. What matters for the threshold router is whether *high-confidence* little predictions are correct on the samples that actually exit. Figure 9 shows this relationship.

The threshold $\tau = 0.70$ is one operating point on a continuum. Figure 10 sweeps τ from 0.50 to 0.95 for C1-C4. Lower thresholds produce cheaper, more exit-heavy cascades; higher thresholds approach big-model behaviour. The cascade-aware variants (C3 and C4) were trained with $\tau = 0.70$ embedded in the soft routing term, and remain competitive

across mid thresholds, but lose their edge once τ is pushed well above the training value.

D. Pico 2 Deployment Results

To check whether the analytical cascade metrics transfer to an embedded target, the selected A2 C1-C4 cascades were exported from PyTorch checkpoints to ONNX and float32 TensorFlow Lite Micro models, then compiled into Raspberry Pi Pico 2 firmware. The independent NAS genotype 4 baseline was exported through the same path. The normal deployment test streamed the full CIFAR-10 test set over USB serial, so each operating point was evaluated on 10,000 images. All Pico runs used the same confidence threshold $\tau = 0.70$ as the main comparison and the same softmax defined in Section II-A.

The A2 variants have almost identical per-path latency because they share the same exported architecture and differ mainly in trained weights and routing behaviour. Their little path takes about 40.3 ms, while a deferred sample takes about 204.2 ms end-to-end. The average latency therefore follows the observed route split: C1 and C3 are the fastest A2 variants because they exit on about 43% of the test set, whereas C4 reaches the highest Pico accuracy but defers 72.4% of samples and is correspondingly slower. The independent baseline exits more often than any A2 variant, but its little and deferred paths are much slower, giving a much higher average latency despite its higher exit ratio.

The deployment also exposed a memory distinction that is not visible from accuracy and MACs alone. The A2 firmware




Fig. 7. Final retrained high-accuracy joint cascades compared with independently searched cascades at the fixed operating threshold $\tau = 0.70$. The plot uses the C4 retraining condition to show the high-accuracy end of the joint ablation, not to define C4 as the single canonical joint result. The left panel compares combined parameter count against cascade accuracy, while the right panel compares expected cascade MACs against cascade accuracy.

TABLE VIII

PICO 2 SERIAL CIFAR-10 DEPLOYMENT RESULTS AT $\tau = 0.70$. **EXIT** AND **DEFER** REPORT THE PERCENTAGE OF THE 10,000 TEST SAMPLES THAT THE THRESHOLD ROUTER EXITED AT THE LITTLE MODEL AND DEFERRED TO THE BIG MODEL RESPECTIVELY. **LITTLE MS** IS THE END-TO-END LATENCY FOR SAMPLES THAT EXIT AFTER THE LITTLE MODEL; **DEFER MS** INCLUDES THE LITTLE PLUS THE BIG MODEL.

Cascade	Acc. (%)	Exit (%)	Defer (%)	Little (ms)	Defer (ms)	Avg (ms)
A2 C1	78.97	42.86	57.14	40.34	204.19	133.97
A2 C2	79.96	32.66	67.34	40.36	204.21	150.70
A2 C3	78.94	42.64	57.36	40.34	204.20	134.33
A2 C4	80.46	27.60	72.40	40.37	204.22	159.00
Independent 4	73.47	58.99	41.01	131.93	848.45	425.77

used separate TensorFlow Lite Micro arenas of 128 KiB for the little model and 256 KiB for the big model, while the independent firmware required arena tuning and was finally tested with a 96 KiB little arena and a 352 KiB big arena. Additional independent energy builds with 160/288, 192/256, 224/224, and 256/192 KiB arena splits were produced to diagnose feasible allocations. This suggests that practical deployment feasibility depends not only on model weights and analytical MACs, but also on activation memory and TensorFlow Lite Micro arena requirements.

E. Energy Consumption

In addition to latency, a smaller route-aware energy experiment was run to measure whether the deployed cascade behaviour translates into lower energy use. Unlike the normal serial test, this energy run used a 1,000-image CIFAR-10 subset embedded directly in Pico flash. This avoided USB streaming during the measurement window and allowed the Pico to be powered through the Uno/shunt path shown in Figure 6. The Pico raised GP18 only during active cascade

inference, so the Uno measurement excluded most idle time. The route-aware build also used GP19 to indicate whether a sample deferred to the big model.

The joint A2 cascade used 34.3 J over the embedded subset, compared with 44.0 J for the independent baseline. This is a 22.0% reduction in total energy and approximately the same reduction per image. Runtime also decreased from 428.7 s to 330.1 s, a 23.0% reduction. The average measured power is similar between the two runs, and is slightly higher for the joint cascade, so the energy saving is mainly explained by the shorter active inference time rather than by a lower instantaneous power draw.

The route counts make this result especially useful. In the energy run, the joint cascade actually deferred slightly more samples than the independent baseline, but still consumed less total energy. This matches the latency result in Table VIII: the joint A2 paths are substantially faster, so the energy benefit does not come only from exiting more often. It comes from the deployed cost of each path. The independent baseline has a high exit ratio, but its little path and deferred path are both

TABLE IX

ROUTE-AWARE PICO 2 ENERGY MEASUREMENT ON A 1,000-SAMPLE EMBEDDED CIFAR-10 SUBSET. **JOINT A2 (C4)** IS THE HIGH-ACCURACY CASCADE-AWARE + LABEL-SMOOTHING VARIANT OF THE CO-SEARCHED A2 CASCADE USED FOR THE DEPLOYMENT COMPARISON. **EXIT** AND **DEFER** ARE THE PERCENTAGE OF THE 1,000 SAMPLES HANDLED BY THE LITTLE PATH AND THE DEFERRED PATH RESPECTIVELY.

Cascade	Time (s)	Energy (J)	J/sample	Power (W)	Exit (%)	Defer (%)
Independent 4	428.7	44.0	0.044	0.099	58.5	41.5
Joint A2 (C4)	330.1	34.3	0.034	0.105	53.0	47.0




Fig. 8. Retraining-loss ablation for the three selected joint architectures at $\tau = 0.70$, plotted in the accuracy-MACs plane. C1 uses CE, C2 uses CE with label smoothing, C3 uses the cascade-aware loss, and C4 combines cascade-aware loss with label smoothing. C3 marginally improves on the CE operating point of C1; the smoothed variants (C2, C4) sit higher on accuracy but also use more MACs because they defer more often.

expensive on the Pico 2. This reinforces the conclusion that exit ratio, analytical MACs, and deployment memory must be interpreted together when judging whether a cascade is actually efficient on hardware.

VII. DISCUSSION

This section interprets the results for the two research questions and the supporting routing analysis. It focuses on what the experiments show about joint cascade design (RQ1), retraining losses (RQ2), and confidence-based routing, and then summarises the main limitations of the study.

A. Joint Cascade Design (RQ1)

The RQ1 comparison supports designing a big/little cascade as a coupled system rather than as two unrelated classifiers. At $\tau = 0.70$ the co-searched cascades reach 80.69% accuracy at 2.08M MACs versus 74.03% at 2.22M MACs for the independent baseline: the joint pipeline is not merely more accurate but more expensive: it occupies a strictly better region of the accuracy-MACs frontier under the same memory envelope.

The mechanism is visible in the per-role MACs breakdown (Appendix A). In the independent baseline the little model is selected as a standalone classifier within a parameter budget,

so it can still be expensive in MACs if it uses costly convolutions on high-resolution feature maps, so parameter count and MACs are not interchangeable. The joint search instead sees the routing behaviour during evaluation and directly optimises expected cascade MACs and cascade accuracy. This pushes it toward cheaper little branches that act as cascade front-ends rather than standalone classifiers, and partly explains why the independent baseline can have a high exit ratio while still producing a weaker accuracy-MACs trade-off.

B. Effect of Retraining Loss (RQ2)

The C1-C4 ablation shows that the retraining objective changes how the selected cascades route inputs, but the NAS-selected architecture remains the dominant factor in the final accuracy-MACs trade-off. CE (C1) is a strong baseline on CIFAR-10. The cascade-aware loss without smoothing (C3) moves in the intended direction, with marginally higher accuracy at slightly lower MACs than C1, but the effect is small. The clearer effect comes from label smoothing: C2 and C4 reduce routing errors and obtain higher accuracy by deferring more samples. C4 (cascade-aware + smoothing) gives the highest mean accuracy and the lowest routing error among C1-C4 but also uses the most MACs. The win is a trade-off, not a strict replacement of CE.

The threshold sweep in Figure 10 adds context: the cascade-aware variants were trained with $\tau = 0.70$ embedded in the soft routing term, and remain competitive at mid thresholds but lose their edge once τ moves well above the training value. The selected architecture, the operating threshold, the confidence behaviour, and label smoothing all shape the final cascade. The cascade-aware loss is best read as a routing-aware regulariser rather than a drop-in CE replacement.

C. Pico 2 Deployment Interpretation

The Pico 2 measurements support the main dynamic-inference assumption behind the analytical MACs metric: early exits are cheaper than deferrals. For the A2 cascades, samples that exit after the little model take about 40 ms, whereas deferred samples take about 204 ms because the big model is evaluated after the little model. This approximately five-fold per-sample latency difference means that the exit ratio has a direct hardware effect. C1 and C3, which exit on about 43% of samples, have lower average latency than C2 and C4, which defer more often. Thus, the measured latency confirms that avoiding the big model produces the expected compute saving on the Pico 2.




Fig. 9. Routing and confidence behaviour at $\tau = 0.70$. *Left*: little-model ECE and routing error for C1-C4. *Right*: cascade accuracy vs. early-exit ratio for joint and independent cascades. Exit ratio alone does not predict cascade quality: it must be read alongside routing error and confidence calibration.




Fig. 10. Threshold sweep for C1-C4. Increasing τ reduces the exit ratio and generally improves accuracy by deferring more samples to the big model.




Fig. 11. Normalized route-aware Pico 2 energy-test comparison. Independent 4 is shown as the 100% baseline; the joint A2 cascade reduces total energy by 22.0% and runtime by 23.0% on the 1,000-sample embedded subset.

At the same time, the hardware results show why analytical MACs should be treated as a proxy rather than a complete deployment model. Within the A2 family, the average analytical

MACs increase from roughly 0.87M MACs for C1 and C3 to roughly 1.03M MACs for C4, and measured latency increases in the same direction. This agreement is useful because it validates the MACs objective as a reasonable ranking signal among closely related cascades. However, the independent baseline breaks any simple MACs-to-latency interpretation. Its average cascade cost is about 2.05M MACs, but its measured latency is 425.77 ms, far more than a simple doubling of the A2 latency would suggest. The independent little model alone takes about 131.93 ms, more than three times the A2 little path, even though both are small enough to fit the search budget.

This mismatch is likely caused by hardware details that are intentionally abstracted away during NAS. The analytical MACs calculation counts multiply-accumulate work, but it does not model TensorFlow Lite Micro kernel implementations, depthwise versus standard convolution efficiency on the Pico 2, memory movement, cache effects, tensor layout conversions, or arena allocation pressure. It also does not

capture the fact that deployment memory is split between flash-resident model data and SRAM-resident tensor arenas. The A2 models used a stable 128/256 KiB little/big arena split, whereas the independent baseline required a much larger big-model arena and several diagnostic builds. This reinforces the distinction made in the methodology: analytical MACs and estimated memory are useful for search and comparison, but final deployment still requires hardware measurement.

The Pico 2 results strengthen RQ1 from a hardware angle. The joint A2 cascades look better on the analytical frontier and also run substantially faster on the target microcontroller. The measured route split explains most of the latency difference within a fixed architecture, while the independent baseline shows that a high exit ratio does not save compute when the little path is itself slow and memory-hungry. MACs, exit ratio, and arena memory must be read together rather than as independent deployment predictors.

D. Limitations

The reported results should be read as the outcome of a practical NAS budget on CIFAR-10, not an exhaustive optimum. The main limitations are:

1) *Proxy Training*: Architecture selection uses a 5-epoch proxy: candidates are judged before they reach full potential. Architectures that learn slowly may be unfairly filtered out, and the final 50-epoch retraining of the survivors does not recover those.

2) *Dataset Size and Complexity*: CIFAR-10 supports the full pipeline at moderate compute, but its limited visual complexity compresses the natural gap between easy little-model samples and hard big-model samples. In several runs, the little model approaches the big model’s standalone accuracy, reducing the room for a cascade-aware loss to exploit a clear role separation. The learned cascade behaviour may therefore depend on CIFAR-10-specific properties (class set, resolution, easy/hard distribution). A more complex dataset could change which architectures are feasible, how often the little model can safely exit, and how large the loss-ablation effects become.

3) *Search Budget*: Joint NAS uses 20 individuals for 20 generations. The independent baseline uses 20 individuals for 10 generations per role. The independent baseline is matched on *number of proxy-trained candidates* (the dominant compute cost), not on evolutionary depth per role. This is a practical compromise that should be tightened in future work. Across the independent baseline the two role-specific searches together consumed approximately 14.4 hours of wall-time on the development machine (summed `time_s` column of `big_search_log.csv` and `little_search_log.csv`); joint-search wall-time was not logged at the per-candidate level and is therefore not directly comparable here, a gap we revisit in future work. Final retraining uses 50 epochs against the 100-200 epochs common in CIFAR-style work, which limits absolute accuracy claims.

4) *Float32 Deployment*: TensorFlow Lite Micro models are deployed at float32, which keeps the export path close to the trained PyTorch models but is not the most memory-efficient

MCU representation. int8 quantisation would free arena budget for larger architectures but would also alter confidence values (and therefore exit decisions); we leave this for future work.

5) *Per-Variant Pico Measurements*: The full 10K serial latency evaluation covers all five operating points (C1-C4 plus the independent baseline), but the route-aware energy experiment was run only on the high-accuracy A2 C4 joint cascade and the independent baseline on a 1K-image embedded subset. Repeating the energy measurement across all C1-C4 variants, on multiple boards, and with quantised models is left to future work.

E. Future Work

Several directions follow directly from the limitations. The most informative next step is to repeat the joint-vs-independent comparison on a more complex dataset (CIFAR-100, Tiny ImageNet) and with a larger search and retraining budget, including a wall-time-matched joint vs independent comparison rather than only a candidate-count match. Further work should sweep the cascade-aware loss hyperparameters ($k, \lambda_l, \lambda_b, \lambda_d, \lambda_w$) across multiple seeds, evaluate alternative routing rules (margin-based, entropy-based, or a small learned router [?], [?]), and extend the deployment study to int8-quantised TFLM models with per-variant energy measurements on multiple boards. Together these would address the proxy, budget, calibration, and hardware-fidelity caveats identified above.

VIII. CONCLUSION

We studied NAS-driven big/little cascades for input-adaptive inference under a 450 KiB MCU deployment budget. The proposed framework combines joint NSGA-II cascade architecture search, constrained model selection, four retraining ablations (C1: CE, C2: CE + smoothing, C3: cascade-aware loss, C4: cascade-aware + smoothing), threshold-based evaluation, and Raspberry Pi Pico 2 deployment. The joint co-search produces substantially better cascades than an independently-searched baseline: at $\tau = 0.70$ it reaches 80.69% accuracy at 2.08 M MACs against the baseline’s 74.03% at 2.22 M MACs, and on-device it cuts runtime by 23% and route-aware energy by 22%. The C1-C4 ablation is more nuanced. C3 nudges accuracy up and MACs down relative to CE but the effect is small. Label smoothing (C2, C4) is the dominant retraining-loss factor. The cascade-aware loss is best read as a routing-aware regulariser whose benefit depends on confidence calibration, exit ratio, and the selected architecture. The Pico 2 results also show that analytical MACs, although a useful search proxy, do not fully predict TensorFlow Lite Micro latency or arena pressure, so MACs, exit ratio, and memory must be interpreted together. Together these results support designing the little model, big model, and routing rule as a single coupled system, and they leave open further work on loss design, calibration, larger datasets, quantisation, and hardware-aware optimisation.

ACKNOWLEDGMENTS

The authors would like to thank Emil Jørgensen Njor and Francesco Daghero for their supervision, feedback, and guidance throughout the project. We also thank the University of Southern Denmark for providing access to computing resources and development facilities used during the experiments.

APPENDIX ADDITIONAL DIAGNOSTIC PLOTS

This appendix collects supporting plots that complement the main RQ1-RQ2 results. They are diagnostic evidence rather than additional headline results.




Fig. 12. Proxy Pareto frontier used to select joint architectures for full retraining. The selected candidates are chosen from the compute-saving knee region after feasibility checks and the constraint $F_{\text{cascade}} < F_b$, not simply by taking the three highest proxy-accuracy points. The figure documents the selection heuristic; final claims are based on fully retrained test-set results rather than the 5-epoch proxy scores.




Fig. 13. MACs breakdown at $\tau = 0.70$. The independent baseline and joint variants use different selected architecture pairs, while the C1-C4 joint ablations use the same selected architectures. Changes in expected cascade MACs within the joint ablation mainly reflect differences in exit ratio caused by the retraining objective.




Fig. 14. Independent-baseline routing diagnostics at $\tau = 0.70$. The independent baseline can have lower ECE than some joint variants but still has lower cascade accuracy, showing that calibration alone is not sufficient to explain the final trade-off.

Confidence vs Accuracy Reliability Diagrams




Fig. 15. Confidence-accuracy reliability diagrams for the C1-C4 retraining variants. The figure complements the ECE and routing-error summaries by showing where each variant is under- or over-confident across confidence bins.